

# Shared memory and Atomics (proposal)

Ecma TC39 January 2016

[lars-t-hansen.github.io/ecmascript\\_sharedmem](https://lars-t-hansen.github.io/ecmascript_sharedmem)

[github.com/lars-t-hansen/ecmascript\\_sharedmem](https://github.com/lars-t-hansen/ecmascript_sharedmem)

# Outline

Motivation for proposal

Technical recap -- if desired

Concerns raised at September '15 meeting

Discussion of proposal's prospects & status

# Motivation

# Goal: Performance

Effective use of multiple cores in ES

2-8 cores are common now

Native apps: threaded + shared memory

ES: one core/worker + slow events

# Use cases (1)

## asm.js and wasm

- “ES as a compiler target”
- Threads and shared memory in translated code
- Compatible with C / C++ / pthreads
- Fast interoperation with ES

In truth, the primary use case.

# Use cases (2)

## “Plain” ES

- Multicore computation with shared state
- Fast communication through shared memory
- Decoupling computation and rendering
- Management of shared assets
- Examples: pdf.js, game AI

# Compromises and constraints

Use cases conflict:

- asm.js has flat memory, no GC, strong types
- “Plain” ES is object-based, GCd, weak types

Can a compromise solution serve both?

For asm.js and wasm, the memory model must be largely compatible with C and C++

# Approach (1)

## Provide low-level facilities

- SharedArrayBuffer + TypedArray
- Atomic operations
- Concurrent agents (workers)
- Agent sleep/wakeup operations



# Approach (2)

Build higher level facilities for ES

- Locks, barriers, synchronic objects
- Communication channels
- Parallel computation abstractions

Extensible Web approach works well here

# Why not only for wasm? (1)

Some tasks are best / only done in ES callouts

- Input and output, DOM access
- “Runtime” tasks

Copy-in / copy-out at boundary too slow (WebGL)

Thus ES wants some shmem access

Also asm.js (= ES) can't polyfill for wasm

# Why not only for wasm? (2)

“wasm-only” does not remove complexity

- ES can call wasm
- wasm will export accessor functions to its shared memory

Now we have threads-as-library: not better

# Why not only for wasm? (3)

Multicore computation in ES is valuable

ES is the better language (GC, rich types)

Divide responsibilities:

- ES code and objects for logic
- Shared memory for shared state
- Shared memory for fast communication
- TypedObjects to structure shared state

# Why not something completely different?

“Compiler target” use case is constraining

Requires shared, racy memory + atomics

“Something different” would be ES-only, not useful as compiler target

Not what this proposal is about

# Status

Spec is stable, change is incremental

Memory model “ok” but not quite done

In Firefox (Nightly) and Chrome (release) now

Apple, Microsoft aware of work, no public signals

API

# API: Shared memory

A new data type:

```
var sab = new SharedArrayBuffer(size)
```

Like ArrayBuffer but:

- cannot be neutered
- memory becomes shared when it is transferred to another agent
- implementations must GC these across agents



# API: Views on shared memory

Existing TypedArrays can be applied to SAB:

```
var sab = new SharedArrayBuffer(size)
```

```
var ia = new Int32Array(sab)
```

```
var fa = new Float64Array(sab, 8, 10)
```

Complicates TA spec a little

Memory access through views is racy, but safe

# API: Atomic operations

## Suite of atomic operations on integer TypedArray

```
Atomics.load(ia, n)
```

```
Atomics.store(ia, n, value)
```

```
Atomics.compareExchange(ia, n, expectVal,  
                        replaceVal)
```

```
Atomics.exchange(ia, n, value)
```

```
Atomics.add(ia, n, value)
```

```
...
```

(All are sequentially consistent; details later)

# API: Agent sleep and wakeup

Modeled on Linux “futex” (fast user-space mutex)

```
Atomics.futexWait(i32a, loc, expect, timeout)
```

```
Atomics.futexWake(i32a, loc, count)
```

Minimal assumptions, very flexible

Fairly hard to use directly

Different use cases will wrap futex differently

# Example: mutex lock()

// State 0 = unlocked, 1 = uncontended, 2 = maybe contended

```
Lock.prototype.lock = function () {  
    const iab = this.iab;           // Int32Array (shared)  
    const idx = this.stateIdx;      // Index in iab  
    var c;  
    if((c = Atomics.compareExchange(iab, idx, 0, 1)) != 0) {  
        do {  
            if(c == 2 ||  
                Atomics.compareExchange(iab, idx, 1, 2) != 0)  
            {  
                Atomics.futexWait(iab, idx, 2);  
            }  
        } while((c=Atomics.compareExchange(iab,idx, 0, 2))!=0);  
    }  
}
```

# Example: mutex unlock()

// State 0 = unlocked, 1 = uncontended, 2 = maybe contended

```
Lock.prototype.unlock = function () {  
    const iab = this.iab;           // Int32Array (shared)  
    const idx = this.stateIdx;      // Index in iab  
    var v0 = Atomics.sub(iab, idx, 1);  
    if (v0 != 1) {  
        Atomics.store(iab, idx, 0);  
        Atomics.futexWake(iab, idx, 1);  
    }  
}
```

# Agent model

Need a model for concurrency in ECMAScript

Define concurrency in terms of agents

Define agents in terms of ES6 jobs

Give jobs a forward progress guarantee

Creating agents, sharing a SAB among agents:  
mapping-specific.

# Agent mapping

In a browser, an agent could be a web worker

- SAB sharing is by `postMessage`
- Web worker semantics need work...
- No `futexWait` on the browser's main thread

# Memory model (1)

Atomics in the program are totally ordered.

Conventional happens-before relation on events:

- Program order (intra-agent)
- Atomic-write  $\rightarrow$  atomic-read (inter-agent)
- `futexWake` called  $\rightarrow$  `futexWait` returns (ditto)
- `postMessage`  $\rightarrow$  event callback (ditto)
- transitivity, irreflexivity



# Memory model (2)

Reads only see writes that happened before them  
(and only the last of those writes)

Unordered accesses, where at least one is a write, create a data race.

Data-race-free programs are sequentially consistent.

# Memory model (3)

Races are safe: Programs don't blow up.

But races are unpredictable: A race poisons memory by reading and writing garbage.

Type-safe values are read (no pointers).

A race affects the union of the locations in the racing accesses.

# Memory model (4)

## Complications:

- Aliased arrays and cells that are not exclusively atomic
- Weakly ordered memory (ARM, MIPS, Power)
- Races have “a little” meaning

## Non-complications:

- No out-of-thin-air values

# Other memory model issues

Only sequentially consistent atomics, for now

- No “relaxed” atomics – we want them, but they are complicated
- No “acquire-release” atomics – ditto

# Concerns Raised

# Concerns

futexWait on the main thread:

- can we avoid it, do we need more features?

Side-channel attacks

- we can build a really good clock

Memory model leaks into the language

- races and optimizations become visible

# futexWait on the main thread

No longer available in Firefox

We get along without it

Have built “asymmetric” data structures to cope:

- futexWait in workers
- Message event on the main thread

Barriers, synchronics, latches, queues

# futexWaitCallback?

Using a Message event is non-compositional

What about this as a better abstraction?

```
// Main thread
```

```
futexWaitCallback(mem, idx, callback)
```

```
// Worker
```

```
futexWake(mem, idx)
```



# futexWaitCallback premature

Better to experiment for a while IMO

Might be useful for “synchronic” primitive

To be investigated further

# Side-channel attack

Can build a good clock in shared memory:

```
// Clock producer thread  
while (true) Atomics.add(mem, idx, 1)
```

```
// Clock consumer thread  
let t = Atomics.load(mem, idx)
```

4ns resolution, fairly reliable

# Cache attack

Clock can be used to launch cache attacks

- cache sniffing (user interactions, secret data)
- row hammering (unclear impact in js)

A clock is needed to distinguish ops that hit or miss the cache

This clock is not needed, but it is better

# Not a new attack in browsers

The cache attack can be launched in other ways

- Flash Player
- Java
- PNaCl/NaCl
- Any native browser extension
  - Chrome native messaging
  - Firefox js-ctypes
  - ActiveX
- (WebAssembly)

# Status quo?

A new capability for JS, not really a new capability in most browsers.

For example, Flash Player:

- Runs automatically for most people
- The rest of the people click “yes” when prompted (says telemetry)

(And then, WebAssembly)

# Mitigations for side-channel?

Mitigation options are so-so:

- Thread affinity solutions are plausible
- Opt-in for shared memory seems difficult

But again, not a new attack in browsers

# Complexity

Races create indeterminacy

ES implementation must not create new races

Atomic operators order operations

ES implementations must not undo that order

Affects semantics and compilers in subtle ways

# Compiler can't introduce races

Must assume unknown function has an atomic op

- Can't move or duplicate TA access across call unless access is known not to be shared

Must honor datum size

- Can't RMW larger datum

Must avoid speculative accesses

- `if (cond) x++;` vs `x++; if (!cond) x--;`



# “Shared memory”

“Shared memory” is independent of the SAB

Atomic ops on one SAB orders all atomic ops

JIT must assume there are unknown SABs

JIT must assume a TA is shared unless proven otherwise

# User affected “only” by races

Data races expose compiler optimizations

```
// One thread (racy load)
while (!mem[flagLoc]) {}
```

```
// Another thread (racy store)
mem[flagLoc] = 1
```

Some JITs hoist the load...

Loops forever if jitted, terminates if interpreted

Proposal's prospects

# Where do we stand? (1)

Does TC39 want this proposal?

Racy memory, true concurrency not negotiable

- high-res timer issue
- spec complexity

Anything else might be negotiated

# Where do we stand? (2)

If TC39 does not take it -- then what?

- competitive issues (native apps more capable)
- web ecosystem (wasm needs something)

Wasm will do its own thing

Wasm programs will just export `peek()`, `poke()`

# Where do we stand? (3)

Ready for stage 2? (“Do we think we want it?”)

- More guidance needed
- Committee involvement (reviewers) desired
- Wasm involvement

(Sundry other slides)

# What's wrong with workers?

- No way to detect if a worker was
  - Started (resource constraints)
  - Terminated
- No forward progress guarantee
- Browser may kill a worker at any time for any reason, clearly overbroad



# Weaker memory models

Acquire/release desirable for performance

```
Atomics.storeRelease(iab, loc, val)
```

```
Atomics.loadAcquire(iab, loc)
```

Relaxed desirable for special cases

```
Atomics.loadRelaxed(iab, loc)
```

```
Atomics.storeRelaxed(iab, loc, val)
```

Excluded to control (initial) complexity

# “Racy-jQuery” (1)

Fear that the web will become racy when a horde of casual scripters start using shared memory

Not unwarranted but overblown IMO

Experiments are probably needed to settle this (ie, deployment and PR)

# “Racy-jQuery” (2)

Restrictions have been suggested:

- Main thread can access shared memory when workers are stopped (callback, probe)
- Main thread has no access to shared memory at all (proxy through workers)
- Workers must opt in to use shared memory (pragma, constructor flag, special module)

But how effective?

# Why not PJS?

PJS subset is hard to pin down

- hard to know what's effective
- impediment to portability, standardization

“Casual” API does not provide speedup

- overhead (recompile, gc, scheduling)
- granularity issues, lack of control

Fixable? Maybe, but not obvious

# Alternate designs (1)

Proper atomic cells or mutexes?

Hard to reconcile with asm.js needs

- Atomic objects require some type of storage management and initialization
- Problematic for pthreads, C++11 `atomic<>`, as well as C11 atomics
- “Rich” features may serve fewer use cases

# Alternate designs (2)

Avoid TypedArray, do something else:

```
let s = new Shared(byteSize)
let v = Shared.int32.load(s, 37)
Shared.int32.store(s, 37, x)
```

Somewhat verbose, user-hostile?

Does not get rid of essential complexity: races